

Remo Sasso

Research scientist in Reinforcement Learning and foundation models, with a PhD in Artificial Intelligence focused on sample-efficient exploration, deep model-based RL, posterior sampling, and foundation model priors for decision-making. Published at **ICML** and **TMLR**, with additional work presented at **NeurIPS workshops**. Currently an **Applied Scientist II at Amazon**, working on personalization, recommender systems, LLM fine-tuning, and transformer-based modelling.

Professional Experience

Amazon

Applied Scientist II

United Kingdom

Oct 2025 – Present

- Work on large-scale applied ML systems across personalization, recommendation, and LLM fine-tuning.
- Internal LLM fine-tuning for product-domain adaptation, including extending model vocabulary for products.
- Lead modelling work on a new requisition-level evaluation model using modern transformer architectures.
- Own experiments end-to-end: problem formulation, implementation, baselines, ablations, evaluation, debugging, failure analysis, and communication of results.

Amazon

Applied Scientist Intern

United Kingdom

Feb 2025 – Jun 2025

- Completed applied research projects at the intersection of personalization, recommender systems, LLMs, and reinforcement learning.
- Designed, implemented, and evaluated experimental ML systems in a production-adjacent research environment.
- Contributed to the launch of a recommendation solution with production impact.

Aletheia, xDNA

Lead AI Developer

Netherlands / Remote

2023 – 2025

- Lead developer of [Aletheia](#), an agentic LLM-based information retrieval and fact-checking platform developed.
- Built systems for retrieval, source grounding, claim analysis, evidence synthesis, structured report generation.
- Panelist at the European Digital Media Observatory Conference 2024 on “AI as a solution for misinformation”.

AIGoalie

Founder

United Kingdom

2024 – Present

- Built a data-driven football forecasting platform using predictive models, automated pipelines, and large-scale fixture data, grown to over 20,000 monthly users.

Selected Publications

Posterior Sampling for Deep Reinforcement Learning

Remo Sasso, Michelangelo Conserva, Paulo Rauber

ICML 2023

- Introduced posterior sampling methods for deep RL to improve principled, sample-efficient exploration.
- [Paper](#) [Code](#)

Multi-Source Transfer Learning for Deep Model-Based Reinforcement Learning

Remo Sasso, Matthia Sabatelli, Marco Wiering

TMLR 2023

- Studied transfer learning across source environments for deep model-based RL and sample-efficient dynamics learning.
- Presented at the **NeurIPS 2022 Deep Reinforcement Learning Workshop**.
- [Paper](#) [Code](#)

Foundation Models as World Models: A Foundational Study in Text-Based GridWorlds

Remo Sasso, Michelangelo Conserva, Dominik Jeurissen, Paulo Rauber

CoG 2026

- Studied whether foundation models can act as world models in text-based reinforcement learning environments.
- Evaluated model-based reasoning, transition prediction, and planning-relevant representations in controlled GridWorld settings.
- [Paper](#)

Education

Queen Mary University of London

London, UK

PhD in Artificial Intelligence

2021 – 2025

Thesis focus: efficient exploration in deep reinforcement learning and foundation model priors for decision-making.

University of Groningen

Groningen, Netherlands

MSc in Artificial Intelligence

2018 – 2021

University of Groningen

Groningen, Netherlands

BSc in Artificial Intelligence

2015 – 2018

Awards & Grants

- **EPSRC IGGI PhD Studentship**, Intelligent Games and Game Intelligence, 2021.
- **Gulbenkian / European Media and Information Fund Grant**, “Boosting Fact-Checking Activities in Europe”, 2023.

Skills

Research: Reinforcement learning, deep RL, model-based RL, exploration, posterior sampling, foundation models, LLMs, VLMs, recommender systems, evaluation, benchmarking, ablations.

Programming: Python, PyTorch, TensorFlow, Scikit-learn, NumPy, Pandas, Matplotlib, SQL, Spark, Java, C, C++, JavaScript, React, Flask, Git, LaTeX.

Infrastructure: AWS, Google Cloud, experiment tracking, data pipelines, model evaluation workflows, production-adjacent ML systems.

Languages: English, Dutch, German, Frysk.